

Christos Hadjinikolis

Machine Learning Engineering Manager | Production ML & Applied AI Systems
Ph.D. Computer Science | UCL Research Associate | AI Standards Expert @ ISO/CEN-CENELEC JTC'21
✉ christos.hadjinikolis@gmail.com

about me

profiles

website
LinkedIn
GitHub
Google Scholar

leadership

6 direct reports
performance/career
delivery ownership
hiring/mentoring

applied AI

LLM/tool use
model eval/replay
streaming UX
approval gates
AI governance

AI stack

PyTorch
Hugging Face
MLflow
SentencePiece
faster-whisper/Kokoro

engineering

model-serving
Python/Java
Kotlin/SQL
Kafka/Flink
reliability

open source

PyPI: dynamicio
skeleton-replay
JetBrains plugin

I am an **ML Engineering Manager** with 16 years across applied ML systems, data platforms, backend services, and AI research. I lead cross-functional teams that translate research-grade models and ambiguous product requirements into reliable, high-throughput production systems. Recent work spans model-serving, batch/online evaluation, streaming ML infrastructure, and reliability for high-volume prediction pipelines, backed by a Ph.D. in Computer Science and active AI standards work with ISO/CEN-CENELEC.

work-experience

Vortexa Ltd, London, UK 12/2020–present

Engineering Manager | ML Systems Lead

I own engineering strategy and delivery for a live ML/data estate turning **6M** vessel-position records/hour into production intelligence for **13.5K** monitored vessels. Manage **6** direct reports and lead a **10-person** cross-functional team accountable for model quality, stakeholder alignment, reliability, and delivery maturity.

- **People & Delivery:** Manage **6** direct reports spanning Product, SME analysis, Data Science, and Data Engineering; own performance/career development, team delivery, hiring for **6** roles, and delivery accountability for the wider team.
- **Model Development & MLOps:** Led 0-to-1 research-to-production delivery for destination/ETA sequence and transformer models in **PyTorch**; established MLflow, model/data versioning, automated evaluation gates, replay, monitoring, and model-serving workflows at **6M** events/hour.
- **Evaluation & Iteration:** Built batch/online model-evaluation and replay loops; analysed edge cases and failure modes with domain experts and Product to drive metric-guided model, data, and interface improvements.
- **Product / Stakeholder Alignment:** Ran workshops across Product, SMEs, analysts, and engineers to turn model-quality disputes into shared definitions, measurable objectives, interface/SLA proposals, and OKR-linked workstreams.
- **Platform & Standardisation:** Standardised I/O seams, data contracts, repo archetypes, AWS CodeArtifact publishing, ADRs, dev containers, and local E2E tests (cut a **3h** DAG feedback loop to under **5** min).
- **Architecture & Reliability:** Led **Kafka Streams-to-Flink** as a strategic platform move; maintained signal-quality controls, rollback/fallback paths, shared on-call, and **MTTR** under **30** min.

DataReply, London, UK 04/2016–12/2020

Senior Consultant | Data Scientist / ML Engineer

Joined Data Reply's London spin-off as its first consultant and helped grow it from **three practitioners and two managers to 30+ consultants**. Progressed to Senior Consultant; shaped opportunities, agreed deliverables/timelines, managed expectations, placed consultants, and delivered technical roadmaps across UBS, Vodafone, CNHi, and other enterprise clients.

- **Vodafone Group:** Led multi-workstream GCP data-science platform (based on Kubeflow) delivery; drove requirements, migration planning, CI/CD, workflow automation, ML lifecycle support, and feature-engineering migration.
- **CNHi:** Lead Data Scientist/Scrum Master on predictive maintenance (Azure/Databricks); translated stakeholder needs into DS/DE backlog and guided a **PySpark** team over live telemetry.
- **UBS:** Embedded consultant who grew from Data Scientist into ML Engineer through XP/pairing while delivering graph/process-mining analytics, observability, and real-time **Kafka**/Elasticsearch pipelines.

KCL, UCL, David Game College, UK 2010–2016

Academic Teaching and Technical Instruction

Taught algorithm design, data structures, databases/SQL, and Java/Python during and after doctoral research; developed presentation discipline, learner support, and mentoring habits.

education

2010–2014	Ph.D. in Computer Science <i>Thesis: Persuasion Dialogues & Opponent Modelling.</i> Developed logical inference methods for large Knowledge Graphs , combining symbolic reasoning, graph analytics , and Bayesian techniques (e.g., MCMC). Received Best Poster Award at IJCAI 2013 (413 submissions), EPSRC PhD Scholarship , Outstanding TA Award , and Graduate Certificate in Academic Practice . Access Publications on Google Scholar	King's College London, UK
2004–2010	Diploma (BEng) in Computer Engineering Five-year polytechnic degree in Computer Science and Engineering with a strong focus on mathematics (10 courses), covering theoretical aspects, networking, and telecommunications. Majored in Artificial Intelligence .	University of Thessaly, GR
2002–2004	Military Officer: Second Lieutenant Completed national service as a commissioned Second Lieutenant.	National Guard, CY

selected AI/tooling work

Private	Promet Private applied GenAI project exploring voice, memory, tool-use, and streaming assistant workflows; hands-on work on local runtimes, faster-whisper STT, Kokoro/ Hugging Face assets, tool routing, schema validation, approval gates, durable state, traces, and deterministic replay/evaluation for auditable LLM interactions.	GenAI harness
PyPI	skeleton-replay Python developer tool that turns script/pytest runs into traces, architecture snapshots, workflow evidence, and replayable reports for review, debugging, onboarding, and LLM-assisted code understanding.	Runtime evidence
JetBrains	Skeleton Replay plugin Free PyCharm/IntelliJ plugin bringing Skeleton reports, workflow evidence, and source navigation into the IDE; reflects a bias for developer experience, evidence, and shared understanding.	IDE workflow
PyPI	dynamicio Free public tooling for making I/O seams explicit, validating schema/data expectations, and switching local/dev/prod datasets so pipeline logic can be tested with faster feedback.	Data contracts

committees/affiliations

since 10/2024	University College London, Dept. of Information Studies Research/teach practical AI standardisation: AI Act, auditability, model/data versioning, explainability, and safe adoption.	Associate Researcher
since 01/2021	ISO - International Organisation for Standardisation Member of JTC'21 Working Group 3, developing AI standards aligned with EU policy and international norms.	Committee Expert Member for JTC'21

talks/interviews

2023	Bill Raymonds, Agile in Action Agile Data Science: Sailing Through Data Science: The Agile Journey at Vortexa, Details: YouTube	Podcast Interview
2022	Open Data Science Conference dynamicio , a published PyPI library for abstracting I/O in ML/data workflows, Details: here	Industry Conference Talk
2020	iunera: Big Data Innovators The Agile approach in data science explained by an ML expert, Details: here	Interview Blog Post
2020	Big Data Warsaw Monitoring & Analysing Communication and Trade Events as Graphs, Details: here	Industry Conference Talk
2018	Connected Data London AI tribes and interconnections: what's graph got to do with it?, Details: here	Panelist @ Industry Conference
2018	Minds Mastering Machines A talk on doing Data Science the Agile Way, Details: here	Industry Conference Talk

certificates & awards

2017–2020	Selected ML/cloud credentials AWS ML Specialty; Google Professional Data Engineer; Apache Spark Developer.	AWS, Google Cloud, O'Reilly
2016–2019	Data/graph engineering credentials Process Mining; Graph Analytics; Advanced Elasticsearch; Neo4j Certified Professional.	Coursera, Elastic, Neo4j