

ChristosHadjinikolis

Staff/Director-Level ML Tech Lead | Ph.D.

Committee Expert Member @ JTC'21

✉ christos.hadjinikolis@gmail.com

mediums
🌐 personal website

 **LinkedIn**
 **github**
 **Google Scholar**
 **YouTube**

languages
bilingual
Greek/English

programming
Python, Java, Kotlin
(Kafka Streams), **SQL**,
Cypher

author
pypi:dynamicio

leadership scope
ML roadmap, delivery
leadership, technical
guidance, supervision,
team scaling,
stakeholder
management,
hiring/interviewing,
mentorship, AI
governance, auditability

cloud-platforms
AWS, GCP

technologies & frameworks
MLOps, MCP, agent
harnesses, Kafka, Kafka
Streams, py/Flink,
FastApi, Docker,
MLflow, Airflow, AWS
CodeArtifact,
Elasticsearch

interests
Sports: Tennis, Jogging,
Reading: Comics,
Philosophy, Politics

about me

Staff-level ML Tech Lead operating at Staff/Director scope, with 15+ years across AI research, production ML systems, MLOps, and developer-facing platform work. Lead initiatives spanning roadmap shaping, prioritisation, stakeholder alignment, architecture standards, mentoring, and real-time inference/event-driven systems. Strong record of scaling R&D impact across systems, people, and product quality by raising engineering standards through CI/CD, schema discipline, lineage, auditability, observability, and pragmatic AI governance.

work-experience

Vortexa Ltd, London, UK 12/2020–present

ML Tech Lead (Staff-Level) | Pod Lead

Lead ML-driven product capabilities for maritime forecasting, combining destination prediction, AI signal cleaning, real-time inference, event-driven architecture, backend services, and production rollout discipline. Influence engineering practices across a 70-person R&D department through technical supervision, product ownership, stakeholder management, backlog prioritisation, mentoring, and standardisation.

- **AI Product Delivery:** Led destination, routing, signal-cleaning, and waiting-time model services, translating R&D into observable, maintainable production capabilities and prediction-quality governance.
- **Event-Driven ML Platform:** Drove **Flink**, **Kafka**, and **Kotlin**/Kafka Streams adoption; shifted stream processing toward framework-driven dataflow, reducing cognitive load and Kafka partition-coupled scaling limits.
- **R&D Scalability:** Raised ML maturity across teams through CI/CD, model/data versioning, schema management, data lineage, auditability, end-to-end tests, ADRs, architecture forums, and repeatable releases.
- **Leadership & Mentorship:** Managed a remote pod, onboarded and mentored engineers, supported hiring/interviewing, and ran an event-driven architecture mentoring programme.
- **Platform Reuse & Governance:** Created **Dynamicio**, adopted across 15+ repositories; implemented **AWS CodeArtifact** adoption for safer internal package publication, reuse, and team-level scalability.
- **AI Developer Tooling:** Identified **LLM/vibe-coding review friction** and initiated MCP/harness work for status lookup, alerting, failure tracking, and audit evidence.
- **Roadmap Balance:** Balanced product delivery, technical-debt reduction, regulatory controls, and stakeholder priorities within a shared ML roadmap.

DataReply, London, UK 04/2016–12/2020

Data Scientist | ML Engineer

Progressed from Data Scientist to ML Engineer, leading cross-functional DS/DE delivery across UBS, Vodafone, and CNHi. Focused on ML platformisation, workflow automation, observability, deployment governance, and maintainable enterprise systems.

- **Vodafone Group (09/2019–11/2020, London):**
 - Worked on Vodafone's internal MLOps platform (**Infinity Project**) on GCP using **Kubeflow**, supporting model versioning, deployment, telemetry, experiment workflows, and resource-governed notebooks.
 - Implemented CI/CD, shared feature/data workflows, and monitoring for reproducible production ML.
- **CNHi (05/2019–09/2019, Chicago):**
 - Built **time-series forecasting** models and transitioned statistical prototypes into scalable cloud services for operational alerting.
- **UBS (08/2017–04/2019, London):**
 - Designed real-time **graph analytics** and **process mining** pipelines for organisational insights and decision automation.
 - Developed observability tooling with **Kafka**, **Elasticsearch**, and Python.

Roles in UK Institutions 2010–2016

Assoc. Lect. @GSM | Coding Teacher @DGC | TA @KCL|UCL

Taught computing during doctoral studies, building deep technical communication habits and the ability to make complex computer science and AI ideas operational for others.

- **2012–2016:** Lectured in Computing, MATLAB, Java, Python, HTML, CSS at GSM and David Game College
- **2010–2014:** TA at King's College London (KCL) and UCL, covering Java, Data Structures, Systems Architecture, Semantic Web, SQL, and AI.

education

2010–2014	Ph.D. in Computer Science <i>Thesis: Persuasion Dialogues & Opponent Modelling.</i> Developed logical inference methods for large Knowledge Graphs , combining symbolic reasoning, graph analytics , and Bayesian techniques (e.g., MCMC). Research contributed to scalable AI systems grounded in formal semantics. Received Best Poster Award at IJCAI 2013 (413 submissions). Awarded EPSRC PhD Scholarship , Outstanding TA Award (student-nominated), and earned the Graduate Certificate in Academic Practice . Access Publications on Google Scholar	King's College London, UK
2004–2010	Diploma (BEng) in Computer Engineering Five-year polytechnic degree in Computer Science and Engineering with a strong focus on mathematics (10 courses), covering theoretical aspects, networking, and telecommunications. Majored in Artificial Intelligence .	University of Thessaly, GR
2002–2004	Military Officer: Second Lieutenant Obligatory army service; position acquired through competitive IQ and athletic exams.	National Guard, CY

committees/affiliations

since 10/2024	University College London, Dept. of Information Studies Support AI application, standardisation, and ethics initiatives in industry and academia.	Associate Researcher
since 01/2021	ISO - International Organisation for Standardisation Member of JTC'21 Working Group 3, developing AI standards aligned with EU policies and international norms.	Committee Expert Member for JTC'21

talks/interviews

2023	Bill Raymonds, Agile in Action Agile Data Science: Sailing Through Data Science: The Agile Journey at Vortexa , Details: YouTube	Podcast Interview
2022	Open Data Science Conference Dynamicio , a python library for abstracting your I/O, Details: here	Industry Conference Talk
2020	iunera: Big Data Innovators The Agile approach in data science explained by an ML expert, Details: here	Interview Blog Post
2020	Big Data Warsaw Monitoring and Analysing Communication and Trade Events as Graphs, Details: here	Industry Conference Talk
2018	Connected Data London AI tribes and interconnections: what's graph got to do with it?, Details: here	Panelist @ Industry Conference
2018	Minds Mastering Machines A talk on doing Data Science the Agile Way, Details: here	Industry Conference Talk

certificates & awards

07/2020	AWS Certified Machine Learning - Specialty Validation Number: D0R0WC7CN1VQ13SV	AWS
01/2019	Process Mining License number: AFU2VTP6273V	Coursera Course Certificates
10/2017	Google Certified Professional - Data Engineer Credential ID: 000000400	Google Cloud Certified
07/2017	Advanced Elasticsearch: Data modelling Enrolment Id: 33442	Elastic
06/2016	Graph Analytics for Big Data License number: NZADYAQ7N8FW	Coursera Course Certificates
05/2016	Apache Spark Developer Certification License number: 1.x-0548	O'Reilly Media
04/2016	Neo4j Certified Professional License number: here	Neo Technology, Creators of Neo4j